

Module 4: Randomized Assignment

Contents

1. INTRODUCTION	2
2. ANALYSIS FOR EVALUATING IMPACTS	2
2.1 Regression Analysis to Quantify Impacts.....	2
2.2 T-test Based Analysis.....	3
2.3 A Decision to Make: Regression Models or T-test?	3
3. RANDOMIZATION IN PRACTICE.....	4
3.1 Simple Randomization.....	4
3.2 Stratified or Block Randomization	5
3.3 Clustered Randomization	5
3.4 Testing the Success of Randomization	6
4. BIBLIOGRAPHY/FURTHER READINGS	7

1. INTRODUCTION

In the previous module we covered causal inference and counterfactual analysis, which are two key concepts used to conduct a rigorous impact evaluation. We also described selection bias and omitted variable bias and how randomization mitigates/eliminates this problem. In the next modules we will cover the various methods typically used when conducting an impact evaluation, including experimental and quasi-experimental methodologies.

This module will focus on randomized assignment. We will spend time making sure you understand how to actually analyze the impacts of a program, where to look, and what to look for. We will explain the various methods that one can use to find the impact, comparing when to use one over the other. Finally, we will walk through more advanced program designs where we may need to stratify on some existing variable (e.g. gender, age, occupation) or cluster at a higher level than the individual (e.g. school, market, village). These ways get at a more precise estimate of the impact, but are also relatively trickier to set up and implement.

At the end of this module, you should be able to:

- ✓ Determine how successful the randomization was (or wasn't)
- ✓ Conduct basic data analysis of an RCT
- ✓ Understand how and when to implement a stratified or clustered randomization

2. ANALYSIS FOR EVALUATING IMPACTS

2.1 Regression Analysis to Quantify Impacts

We can write the observed outcome Y for an individual i as,

$$Y_i = T_i \cdot Y_{i_{trt}} + (1 - T_i) \cdot Y_{i_{ctr}}$$

where $T_i = 1$ if the individual i is assigned to a treatment group and $T_i = 0$ if the individual is assigned to the control group, and Y_i is that individual's observed outcome. The *trt* and *ctr* are used to clarify that the individual can be either in treatment or control groups in "real life" analysis. Rearranging the terms

$$Y_i = Y_{i_{ctr}} + (Y_{i_{trt}} - Y_{i_{ctr}}) \cdot T_i, \text{ and}$$

$$Y_i = E[Y_i]_{ctr} + (Y_{i_{trt}} - Y_{i_{ctr}}) \cdot T_i + (Y_{i_{ctr}} - E[Y_i]_{ctr})$$

which in usual linearity regression notation, assuming linearity, can be represented as,

$$Y_i = \beta_0 + \beta_1 \cdot T_i + \varepsilon_i$$

Based on the above regression model, we can estimate the conditional outcome with and without the treatment T and then estimate the causal effect as follows,

$$T_i = 1: \quad E[Y_i|T = 1] = \beta_0 + \beta_1 + \varepsilon_i$$

$$T_i = 0: \quad E[Y_i|T = 0] = \beta_0 + \varepsilon_i$$

$$E[Y_i|T = 1] - E[Y_i|T = 0] = \beta_0 + \beta_1 + \varepsilon_i - (\beta_0 + \varepsilon_i) = \beta_1.$$

Therefore, coefficient β_1 quantifies the impact as group mean difference in outcomes between the treatment and control group. Remember, estimates of β_1 obtained in this way are unbiased only if selection bias is 0.

Exercise 4.1: Open `PanelPROGRESA_97_99year.csv`. This dataset is a repeated cross-section of different waves of the ENCEL survey for March and October 1998 and November 1999. It also includes the baseline data collected in 1997. We have used this or part of this dataset previous modules as well.

Let's assume that we expect PROGRESA (`D_HH`) to change the income of the household (`IncomeLabHH`) in 1999. Specify a regression model in R as discussed above and restrict the data to be used from year 1999 (variable `year`). What is the impact of the intervention on income levels? Is it statistically significant? How do you interpret the coefficient?

2.2 T-test Based Analysis

We have amply practiced t-test to detect difference between the two groups; for example, whether household assets value is different between highly-educated versus poorly-educated household heads (Module 1.3).

Here, we can extend t-test analysis to comparing the two groups which differ by their treatment assignment (T_i). As before, the null hypothesis is: the outcome of treated individuals is the same as it would have been had those individuals not been treated. The alternative hypothesis can be specified as a two-sided or one-sided (one is larger/smaller than other) comparison.

Remember, we are making an assumption that the individuals receiving the treatment are exactly like those not receiving the treatment because of the “missing data” problem. In other words, randomization and independence implies that there is no selection bias.

Exercise 4.2: Conduct a two-sided t-test to compare if household income is different by the treatment assignment. Are the results same as those from the regression analysis?

2.3 A Decision to Make: Regression Models or T-test?

The above exercise demonstrates that causal effects are the same whether you use regression analysis or t-test. However, both of these analyses assumed that selection bias was zero. Was this the case? Whether randomized or non-randomized design is used, it is possible that the comparison groups were imbalanced at the baseline. What if the randomization is stratified by geography, for instance (we will discuss randomization strategies in a later module)? If that was the case, we would have to

Learning Guide: Randomized Assignment

account for how we modeled the individuals in our analysis. What if treatments were conducted at the village level, but we believe that individuals within a village are influenced by each other and share common facilities, so that there is some “correlation” among their behaviors? Then we have to estimate the standard errors clustered at the village level (refer to Module 1.3 to learn more). Also, we will later learn more robust methods, such as difference-in-difference, which provide additional robustness relative to regression analysis in some cases.

In general, impact evaluations often face problems that have to be “controlled for” or adjusted for in the analysis.

Regression methods give us the tools to do such adjustment which a simple t-test cannot. T-tests (or chi-squared tests if the outcome is categorical) alone are valid only if our groups are properly randomized, sample size is adequately large to achieve baseline balance in the two groups, data collection is unbiased, and you can reasonably adhere to the standard assumptions of t-test.

Exercise 4.3: As a demonstration, let’s assume that poverty status of a household (`pov_HH`) is a confounder and we should control for it in the regression analysis. Specify a regression model in R that tests the effect of PROGRESA on income, controlling for poverty status of the household and remembering to restrict the data to be used from year 1999 (variable `year`).

Now, specify a separate regression model, this time clustering the standard errors at the village level. We can most easily accomplish this using the `lm.cluster` function in R:

```
PROGRESA_99 <- PanelPROGRESA_97_99year[which(PanelPROGRESA_97_99year$year
== 1999), ]
model2 <- lm.cluster(data = PROGRESA_99, formula=IncomeLab_HH ~ D_HH+pov_HH,
cluster=PROGRESA_99$villid)
summary(model2)
```

You will need to install and load the packages `miceadds` and `multiwayvcov` to use this command.

What happens to the effect size (magnitude of the coefficient) and the standard errors when comparing output from the two models? Based on the output from these regressions, did PROGRESA significantly impact income levels?

3. RANDOMIZATION IN PRACTICE

Typically, there are important experimental steps to take before you randomize the treatment group. These steps include: (1) based on several assumptions, you estimate the sample size you would need to detect the impact you are interested in (covered in later modules); and (2) you select a population which you can include in your experiment. This can be a purposive selection in discussion with intervention implementing agencies, a representative sample of the target population and anything in between. Here, we discuss how to randomly assign the people into the treatment and control groups. We will cover the quasi-experimental selection of control groups later.

3.1 Simple Randomization

Under a simple randomization framework each individual, household, or any other “unit of analysis”

Learning Guide: Randomized Assignment

has equal chance of being selected in treatment or control groups; that is, a 0.5 probability of being part of each group. You can imagine flipping a coin and assigning the individual to the treatment group if the coin turns heads and to the control group if it is tails. You can conduct such a randomized allocation through R as demonstrated below.

- ✓ Open `PanelPROGRESA_97_99year.csv` if you have not already
- ✓ Generate a random variable which takes value of 0 or 1 as:

```
PanelPROGRESA_97_99year$random_T <-
round(runif(nrow(PanelPROGRESA_97_99year)), 0)
```

- ✓ You must perform randomization before the program or at the baseline (whether or not you conduct a baseline survey). However, in this example, we have data for 1998 and 1999 already so we restricted the randomization to 1997 observations only. Set the other values to NA

```
PanelPROGRESA_97_99year$random_T <-
ifelse(PanelPROGRESA_97_99year$year == 1997,
PanelPROGRESA_97_99year$random_T, NA)
```

- ✓ Note, anytime you use a random generator (e.g. `runif` function above), it is best to set `.seed0` to a fixed value so that you can always reproduce the results. `runif()` creates a uniform distribution which is basically random number generation.

3.2 Stratified or Block Randomization

In some experiments it is practical or theoretically proper to randomize at some unit of grouping (called blocks or strata). For example, you may want to control for regional administrative, ecological, or political factors and maximize “exchangeability” by randomizing at an administrative district or block level. It is possible that you have selected your study participants in a 2-stage sample and now you want to randomize the treatment in the same manner. Or you are concerned with attrition (individuals leaving your sample before you conduct follow-up analysis) so that at the end of the study you will be faced with “un-exchangeable” groups. In this case, you can: (a) break the study sample in random groups of sizes (for example) 2, 4, and 6 villages each. Within each group, you can randomize the treatment – half to treatment and half to controls. This is done so that even if you lose controls in a “group” you can discard that group of 2, 4 or 6 villages but still be assured of exchangeability in the remaining groups. In short, there are strong reasons for stratifying the randomization. It can be conducted in R as follows:

- ✓ Suppose we want to stratify by federal administrative regions (`geopolid`)
- ✓ We will repeat what we did in 3.1 but just stratified as,


```
PanelPROGRESA_97_99year <- PanelPROGRESA_97_99year %>%
group_by(geopolid) %>% mutate(random_strata_T =
round(runif(n()), 0))
```
- ✓ `runif(n())` randomly assigns 0 and 1 within each group whose size is returned by `n()`

3.3 Clustered Randomization

This is a practical and popular strategy often employed in the development sector. Most development interventions or programs are not targeted at individuals but at some cluster: a village, an office department, or some other group of people. All individuals within the targeted cluster can be the intended customers/beneficiaries/target of the intervention, or there can be a “selection criteria” or

eligibility criteria within the cluster, or there may be random selection of participants within the selected cluster for the program (the latter is termed “block randomization”). In R you can do cluster randomization as follows.

- ✓ Create an identifier that flags unique records at a cluster level. For example, suppose we want to randomize the villages then we should identify one unique observation for each village. We can do this as `PanelPROGRESA_97_99year$uniqvill <- PanelPROGRESA_97_99year %>% group_indices(villid)`
- ✓ Now, randomly assign half of the villages to treatment group as follows, `PanelPROGRESA_97_99year <- PanelPROGRESA_97_99year %>% group_by(uniqvill) %>% mutate(random_cluster_T = round(runif(1), 0))`
- ✓ We’re using `runif(1)` here as we need the same value for all rows in the group, but random assignment of 1s and 0s across groups

3.4 Testing the Success of Randomization

We discussed how the randomization assumption is based on large sample size to randomize the treatment so that we can assume that all measured and unmeasured confounders are distributed equally in the comparison groups. However, in reality we have to assess whether randomization was “successful” in achieving balance. The way to test for “success of randomization” is by evaluating the balance in measured variables between the treatment and control groups at the baseline, which should be the first table in any kind of report. Consider the following:

- ✓ We have to use baseline data because some of the measured variables can be affected by the intervention, and they can change differentially in the treatment and control group, so that comparing them after the commencement of a treatment might be biased
- ✓ Just because we find statistically significant difference at the baseline does not mean that the groups are imbalanced. For example, continuous measurements of age (years) and income (US Dollars) will likely be statistically different between the two groups even if the sample size of each is very large, because large sample size allows us to detect even very small differences in continuous variables. Therefore, we should assess only whether the difference is economically, biologically and logically large, not necessarily relying only on statistical significance.
- ✓ The converse of above is also true. Just because the difference is statistically insignificant does not mean that the groups are well balanced, if the differences have large magnitudes.
- ✓ It is best to compare the treatment and control groups on all available measurements for balance at the baseline so that you can be reasonably confident that the comparison groups are balanced, at least on observables. Best practice is to select these comparison variables “before” you randomize and faithfully check the group mean difference after the randomization at baseline. Plus, we can always add this covariate into our regression afterwards to check for any potential omitted variable bias!

For the dataset example we have been following so far, let’s demonstrate how to check for the balance.

- ✓ Run the following commands to check the balance for a few selected variables at the baseline.

Note, D indicates whether the village was randomized to treatment group or not.

```
t.test(IncomeLab_HH ~ D, data = PanelPROGRESA_97_99year[which(PanelPROGRESA_97_99year$year == 1997), ], var.equal = TRUE, alternative = "two.sided")
t.test(famsize ~ D, data = PanelPROGRESA_97_99year[which(PanelPROGRESA_97_99year$year == 1997), ], var.equal = TRUE, alternative = "two.sided")
t.test(eduhead ~ D, data = PanelPROGRESA_97_99year[which(PanelPROGRESA_97_99year$year == 1997), ], var.equal = TRUE, alternative = "two.sided")
t.test(agehead ~ D, data = PanelPROGRESA_97_99year[which(PanelPROGRESA_97_99year$year == 1997), ], var.equal = TRUE, alternative = "two.sided")
```

- ✓ Figure 4 is the output that shows how well the groups are balanced for one factor. You can look at the outputs for each factor and find that all the factors are balanced very well, but some of the differences are statistically significant.
- ✓ In the case of stratified or block randomization, you should evaluate the balance between each block or strata.

```
> t.test(famsize ~ D, data = PanelPROGRESA_97_99year[which(PanelPROGRESA_97_99year$year == 1997), ], var.equal = TRUE, alternative = "two.sided")

Two Sample t-test

data: famsize by D
t = 1.7667, df = 24075, p-value = 0.07729
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.006624202  0.127695696
sample estimates:
mean in group Control mean in group Treated
      5.257022          5.196486
```

Figure 4. Testing for baseline balance

4. BIBLIOGRAPHY/FURTHER READINGS

1. Duflo, Esther; Rachel Glannester and Micheal Kremer (2008). "Using Randomization in Economic Development Research: A Toolkit," *Handbook of Development Economics*, Vol. 4, Elsevier Science.
2. Gerber, Alan S., and Donald P. Green. "Field experiments: Design, analysis, and interpretation." WW Norton, 2012.
3. Gertler, Paul J., Sebastian Martinez, Patrick Premand, Laura B. Rawlings, and Christel MJ Vermeersch. "Impact evaluation in practice." World Bank Publications, 2011.